

slideSonnet

Narration for your slides — as editable text

A self-narrated tour

What is this?

Alex: So what does slideSonnet actually do?

What is this?

Alex: So what does slideSonnet actually do? **Narrator:** You bring a finished PDF of your slides. slideSonnet turns it into a narrated video — and the narration you're hearing right now lives in a plain text file next to that PDF.

Why a sidecar file?

Alex: Why not just put the narration inside the slides?

Why a sidecar file?

Alex: Why not just put the narration inside the slides? **Narrator:** Because the PDF is your frozen artifact. Keeping narration in a separate, git-diffable file means you can rewrite a line and re-render in seconds — no recompiling, no re-recording.

Step 1 — mark your slides

Add the macro to your Beamer source and give every slide a stable id:

```
\usepackage{slidesonnet}  
...  
\ssid<1>{euler-setup}      % overlay step 1  
\ssid<2>{euler-trick}     % overlay step 2  
...
```

The id is stamped invisibly onto each page — narration binds to it, not to a page number.

Step 2 — scaffold the sidecar

One command reads the ids straight out of the PDF:

```
slidesonnet init deck.pdf
```

You get a blank `deck.narration` with one block per slide, ready to fill in — by hand, or by an LLM.

The sidecar format

```
@euler-setup
:voice narrator
We want the sum of one over n squared. [pause 0.8]

@euler-trick
:pace slow
Watch the denominators. [pause 1] This is the trick.
```

@id headers, optional :voice/:pace, and [pause N] — the one timing primitive.

Step 3 — edit and preview

Alex: Do I have to write that text file by hand?

Step 3 — edit and preview

Alex: Do I have to write that text file by hand? **Narrator:** You can — or run `slidesonnet edit deck.pdf` for a local editor. Page through the deck, type narration beside each slide, and preview the whole thing. The preview bakes in your pauses, so it plays exactly like the final video.

Step 4 — render

```
slidesonnet export deck.pdf -o deck.mp4 --engine piper
```

- `--engine piper` — free, local. `elevenlabs` — cloud, high quality.
- `--silent` or `--timing estimate` — a fast rough cut, no TTS.
- Subtitles (`.srt/.vtt`) are written alongside the video.

Only what changed

Alex: Doesn't re-rendering get slow?

Only what changed

Alex: Doesn't re-rendering get slow? **Narrator:** Audio is cached by content. Edit one block and only that line is re-synthesized; everything else is reused. And swapping Piper for ElevenLabs never throws away the audio you already paid for.

The whole workflow

```
slidesonnet sty                # drop the LaTeX macro
slidesonnet init    deck.pdf    # scaffold narration
slidesonnet check   deck.pdf    # reconcile ids
slidesonnet edit    deck.pdf    # write & preview
slidesonnet export  deck.pdf -o deck.mp4
```

Every step is scriptable — an LLM or a Makefile can drive the whole pipeline.

slideSonnet

Mark your slides. Write narration. Render.

```
pip install slidesonnet
```