

SlideSonnet

TEXT → VIDEO

Narration for your slides — as editable text

A self-narrated tour

ALEX

So what does slideSonnet actually do?

ALEX

So what does slideSonnet actually do?

NARRATOR

It turns your slide deck into a narrated lecture video. Voices, page timing, subtitles — all driven by one plain text file.

ALEX

Why not just put the narration inside the slides?

ALEX

Why not just put the narration inside the slides?

NARRATOR

Because the PDF is your frozen artifact. Keeping narration in a separate, git-diffable file means you can rewrite a line and re-render in seconds — no recompiling, no re-recording.

Add the macro to your Beamer source and give every slide a stable id:

```
\usepackage{slidesonnet}  
...  
\ssid<1>{euler-setup}      % overlay step 1  
\ssid<2>{euler-trick}     % overlay step 2  
...
```

The id is stamped invisibly onto each page — narration binds to it, not to a page number.

One command reads the ids straight out of the PDF:

```
slidesonnet init deck.pdf
```

You get a blank `deck.narration` with one block per slide, ready to fill in — by hand, or by an LLM.

```
@format
  utterance:
    text: Each slide id gets a block of utterances and pauses.
  utterance:
    voice: alex
    text: Wait. Is that the block we're hearing right now?
  utterance:
    pace: slow
    text: Yes. Word for word.
```

One block per slide id; every utterance carries its own `voice:` and `pace:` — a single slide can host this whole exchange.

ALEX

Do I have to write that text file by hand?

ALEX

Do I have to write that text file by hand?

NARRATOR

You can — or run `slidesonnet edit deck.pdf` for a local editor. Page through the deck, type narration beside each slide, and preview the whole thing. The preview bakes in your pauses, so it plays exactly like the final video.

```
slidesonnet export deck.pdf -o deck.mp4 --engine kokoro
```

- `--engine kokoro` — free, local, natural. `elevenlabs` — cloud, high quality.
- `--silent` or `--timing estimate` — a fast rough cut, no TTS.
- Subtitles (`.srt/.vtt`) are written alongside the video.

ALEX

Doesn't re-rendering get slow?

ALEX

Doesn't re-rendering get slow?

NARRATOR

Audio is cached by content. Edit one block and only that line is re-synthesized; everything else is reused. And switching engines never throws away audio you already paid for.

```
slidesonnet sty          # drop the LaTeX macro
slidesonnet init  deck.pdf # scaffold narration
slidesonnet check  deck.pdf # reconcile ids
slidesonnet edit   deck.pdf # write & preview
slidesonnet export deck.pdf -o deck.mp4
```

Every step is scriptable — an LLM or a Makefile can drive the whole pipeline.

SlideSonnet

TEXT → VIDEO

Mark your slides. Write narration. Render.

```
uv tool install "slidesonnet[kokoro]"
```