

SlideSonnet

TEXT → VIDEO

Narration for your slides — as editable text

A self-narrated tour

ALEX

So what does slideSonnet actually do?

ALEX

So what does slideSonnet actually do?

NARRATOR

It compiles narration from plain text. Write what each slide should say, a synthesized voice reads it, and your PDF becomes a timed, subtitled lecture video.

ALEX

Why a computer voice? I could record myself.

ALEX

Why a computer voice? I could record myself.

NARRATOR

Because slides change — and every change means re-recording. Text doesn't mind: fix the sentence, re-render, done. Your video's source becomes git-diffable text, versioned like code.

Add the macro to your Beamer source and give every slide a stable id:

```
\usepackage{slidesonnet}  
...  
\ssid<1>{euler-setup}      % overlay step 1  
\ssid<2>{euler-trick}     % overlay step 2  
...
```

The id is stamped invisibly onto each page — narration binds to it, not to a page number.

One command reads the ids straight out of the PDF:

```
slidesonnet init deck.pdf
```

You get a blank `deck.narration` with one block per slide, ready to fill in — by hand, or by an LLM.

```
@format
utterance:
  text: Each slide id gets a block of utterances and pauses.
utterance:
  voice: alex
  text: Wait. Is that the block we're hearing right now?
utterance:
  pace: slow
  text: Yes. Word for word.
```

One block per slide id; every utterance carries its own `voice:` and `pace:` — a single slide can host this whole exchange.

ALEX

Do I have to write that text file by hand?

ALEX

Do I have to write that text file by hand?

NARRATOR

You can — or run `slidesonnet edit deck.pdf` for a local editor. Page through the deck, type narration beside each slide, and preview the whole thing. The preview bakes in your pauses, so it plays exactly like the final video.

```
slidesonnet export deck.pdf -o deck.mp4 --engine kokoro
```

- `--engine kokoro` — free, local, natural. `inworld` — cloud, high quality.
- `--silent` or `--timing estimate` — a fast rough cut, no TTS.
- Subtitles (`.srt/.vtt`) are written alongside the video.

ALEX

Doesn't re-rendering get slow?

ALEX

Doesn't re-rendering get slow?

NARRATOR

Audio is cached by content. Edit one block and only that line is re-synthesized; everything else is reused. And switching engines never throws away audio you already paid for.

```
slidesonnet sty          # drop the LaTeX macro
slidesonnet init    deck.pdf      # scaffold narration
slidesonnet check    deck.pdf      # reconcile ids
slidesonnet edit     deck.pdf      # write & preview
slidesonnet export   deck.pdf -o deck.mp4
```

Every step is scriptable — an LLM or a Makefile can drive the whole pipeline.

SlideSonnet

TEXT → VIDEO

Mark your slides. Write narration. Render.

```
uv tool install "slidesonnet[kokoro]"
```