

# SlideSonnet

TEXT → VIDEO

*Narration for your slides — as editable text*

A self-narrated tour

**Alex:** So what does slideSonnet actually do?

**Alex:** So what does slideSonnet actually do? **Narrator:** You bring a finished PDF of your slides. slideSonnet turns it into a narrated video — and the narration you're hearing right now lives in a plain text file next to that PDF.

**Alex:** Why not just put the narration inside the slides?

**Alex:** Why not just put the narration inside the slides? **Narrator:** Because the PDF is your frozen artifact. Keeping narration in a separate, git-diffable file means you can rewrite a line and re-render in seconds — no recompiling, no re-recording.

Add the macro to your Beamer source and give every slide a stable id:

```
\usepackage{slidesonnet}  
...  
\ssid<1>{euler-setup}      % overlay step 1  
\ssid<2>{euler-trick}     % overlay step 2  
...
```

The id is stamped invisibly onto each page — narration binds to it, not to a page number.

One command reads the ids straight out of the PDF:

```
slidesonnet init deck.pdf
```

You get a blank `deck.narration` with one block per slide, ready to fill in — by hand, or by an LLM.

```
@euler-setup
:voice narrator
We want the sum of one over n squared. [pause 0.8]

@euler-trick
:pace slow
Watch the denominators. [pause 1] This is the trick.
```

@id headers, optional :voice/:pace, and [pause N] — the one timing primitive.



**Alex:** Do I have to write that text file by hand?

**Alex:** Do I have to write that text file by hand? **Narrator:** You can — or run `slidesonnet edit deck.pdf` for a local editor. Page through the deck, type narration beside each slide, and preview the whole thing. The preview bakes in your pauses, so it plays exactly like the final video.

```
slidesonnet export deck.pdf -o deck.mp4 --engine piper
```

- `--engine piper` — free, local. `elevenlabs` — cloud, high quality.
- `--silent` or `--timing estimate` — a fast rough cut, no TTS.
- Subtitles (`.srt/.vtt`) are written alongside the video.

**Alex:** Doesn't re-rendering get slow?

**Alex:** Doesn't re-rendering get slow? **Narrator:** Audio is cached by content. Edit one block and only that line is re-synthesized; everything else is reused. And swapping Piper for ElevenLabs never throws away the audio you already paid for.

```
slidesonnet sty          # drop the LaTeX macro
slidesonnet init    deck.pdf    # scaffold narration
slidesonnet check    deck.pdf    # reconcile ids
slidesonnet edit     deck.pdf    # write & preview
slidesonnet export   deck.pdf -o deck.mp4
```

Every step is scriptable — an LLM or a Makefile can drive the whole pipeline.

# SlideSonnet

TEXT → VIDEO

Mark your slides. Write narration. Render.

```
pip install slidesonnet
```